

**Lab:** Response Time Analysis using FpsCalc  
**Course:** Real-Time Systems  
**Period:** Autumn 2021

## Lab Assistants

Guru Mehar Rachaputi, Renjith Ramachandran, Shenghui Li, and Gaoyang Dai

## Introduction

The purpose of this assignment is to give students a chance to practice the theories covered by lecturers and course literature. The assignment will be focusing on the Rate Monotonic (RM) and Fixed Priority Schedulability (FPS) analyses. The assignment will also cover how to take into account factors like blocking and jitter in the analysis.

The assignment should be solved and submitted by groups. For more information on report procedures please see section Report on page 11.

## FpsCalc

FpsCalc is a tool for performing Fixed Priority Schedulability analysis and is a recommended help for solving the given assignments.

Please read FpsCalc User's Manual downloaded from the lab webpage carefully to get an idea about how to use it. If you are sitting at a Linux-machine then you can download the *fpscalc\_linux.tar.gz* package and extract it in a folder. Use command *cd* to go into the folder *fpscalc\_linux* and execute FpsCalc by typing following command (your program is the name of the file to be calculated):

```
./fpscalc < your_program.fps
```

## Notation

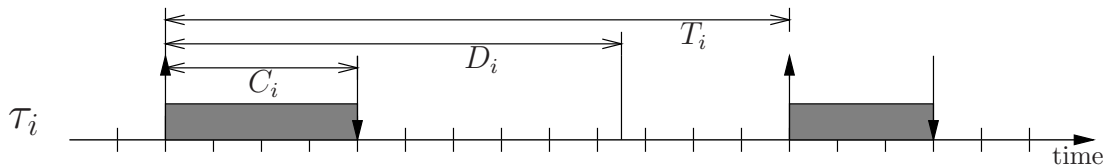


Figure 1: Typical parameters of a real-time task

We can illustrate a tasks timely behavior using a time line as in Figure 1, where the gray regions indicate that the task  $\tau_i$  is executing. We will use the following notation for describing tasks:

$T_i$  is the period of task  $\tau_i$ . For aperiodic tasks we use  $T_i$  to denote the *minimum inter-arrival time* for task instances.

$C_i$  is the Worst Case Execution Time (WCET) of task  $\tau_i$ . Please observe that  $C_i$  should be a safe upper estimate of the execution time for the task, but this does not mean that  $\tau_i$  will execute whole  $C_i$  time every time it gets released (as illustrated in Figure 1 on the previous page).

$D_i$  is the relative deadline of task  $\tau_i$ . I.e. the maximum allowed time between the arrival (release) and the completion of an instance.

$P_i$  is the priority of task  $\tau_i$ .

$B_i$  is the longest time task  $\tau_i$  might be blocked by lower priority tasks, (see Section Blocking on page 5).

$J_i$  is the worst case jitter for task  $\tau_i$  (see Section Jitter on page 8).

$R_i$  is the worst case response time for task  $\tau_i$ . The goal of most schedulability analysis is to find  $R_i$  and verifying that  $R_i \leq D_i$ .

We also let  $lp(i)$ ,  $ep(i)$  and  $hp(i)$  denote the set of tasks with priority less than, equal to and higher than task  $\tau_i$  respectively. In the  $ep(i)$  also task  $\tau_i$  will be included.

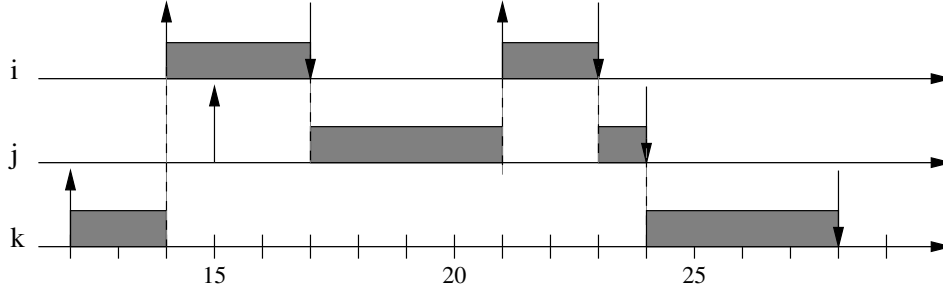


Figure 2: Example of a preemptive schedule

In fixed priority scheduling each task will have a priority assigned to it before execution which will not change over time.

A system which have several tasks executing can be illustrated by a time table as in Figure 2. The switch between tasks can either be *preemptive* - the kernel can stop a task to let another task run, and *non-preemptive* - a task that has started must complete its execution before we can start another task. The example given in Figure 2 illustrates an execution using a preemptive scheduler, where task  $\tau_i$  has the highest and  $\tau_k$  the lowest priority. An  $\uparrow$ -arrow indicates that a task gets released and an  $\downarrow$ -arrow that a task has finished its execution.

In more detail, the table in Figure 2 illustrates one execution where task  $\tau_k$  starts executing at time 12. At time 14 task  $\tau_i$  gets released and preempts task  $\tau_k$ . At time 15 task  $\tau_j$  gets released but can not start executing because the higher priority task  $\tau_i$  is executing. When a task finish executing the highest priority waiting task are allowed to execute again, as happens at time 17 when task  $\tau_j$  starts to execute. At time 21 a new instance of the high priority task  $\tau_i$  gets started and preempts  $\tau_j$ . Finally at time 24 when both task  $\tau_i$  and  $\tau_j$  has finished their executions task  $\tau_k$  can continue its preempted execution.

You should preferably use this type of time tables and the given notations when presenting solutions to the given assignments.

## Rate Monotonic

The *Rate Monotonic* (RM) priority ordering assigns priorities to tasks according to their periods. Specifically, tasks with higher request rates, i.e. tasks with smaller periods, get higher priorities.

I.e. the task with the shortest period gets the highest priority and the task with longest period gets the lowest priority. Tasks with higher priority can preempt lower priority tasks. To get the RM analysis to work all the tasks must have the characteristics of being periodic, independent and having deadline equal to period. The RM priority assignment is *optimal* meaning that if the task set is schedulable with a given fixed-priority assignment it is also schedulable with the RM priority assignment.

The *utilization* of a system of  $n$  tasks is the fraction of total execution time spent in executing the tasks, i.e. the computational load of the system. An upper *utilization bound*  $U$  of a system with  $n$  periodic tasks can be derived as:

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \quad (1)$$

It can be shown that a set of  $n$  independent periodic tasks, with deadlines equal to periods and scheduled by the RM algorithm will always meet its deadlines, for all task phasing, if

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq n(2^{1/n} - 1) \quad (2)$$

We can see that when  $n \rightarrow \infty$  the  $n(2^{1/n} - 1)$  expression approaches  $\ln 2 \approx 0.693$ .

The *critical instant* of a task is the time at which the release of a task will produce the largest response time for the task. There exists a theorem saying that the critical instant of a task  $\tau_i$  occurs whenever the task is released simultaneously with all tasks with higher priority than  $\tau_i$ , (i.e. the tasks in  $hp(i)$ ). The most of the following discussions will rely on this theorem for their correctness.

The simple version of the critical instant theorem is only valid when the tasks are independent, fixed priority and runs with perfect periodicity. In the latter sections you will investigate how to change the theorem and corresponding response time formulas to handle factors like blocking and jitter.

## Assignment 1

Task	$C_i$	$T_i$	$D_i$
$\tau_1$	2 ms	10 ms	10 ms
$\tau_2$	4 ms	15 ms	15 ms
$\tau_3$	10 ms	35 ms	35 ms

Figure 3: A number of periodic tasks with  $D_i = T_i$

Given the task set in Figure 3:

- 1.1 What is the priority ordering for the tasks using the RM priority ordering?
- 1.2 Will all tasks complete before their deadlines according to the schedulability formula in Equation 2 on page 3? Draw a critical instant schedule for the given tasks (as in Figure 2, but have all tasks simultaneously released at time 0). What is the system utilization bound?
- 1.3 Assume that we keep all the parameters in Figure 3 and only increase the computation time for task  $\tau_1$  to be  $C_1 = 5$ . Will all task complete before their deadlines? Draw a critical instant schedule for the given tasks. What is the system utilization bound?
- 1.4 Assume that we instead of modifying  $\tau_1$ , (set  $C_1 = 2$ ), want to increase the computation time for task  $\tau_3$  to be  $C_3 = 17$ . Will all task complete before their deadlines? Draw a critical instant schedule for the given tasks. What is the system utilization bound?
- 1.5 What conclusion can you draw from all this for the schedulability formula in Equation 2 on page 3?

Specify your conclusion in terms of the system utilization bound, the  $n(2^{1/n} - 1)$  expression and 1.00.

- 1.6 Now, assume that we change the relative deadline for task  $\tau_3$  in Figure 3 on the previous page to be  $C_3 = 15$ . Please draw a demand bound function (check lecture slide scheduling theory part2.pdf) of the changed task set and explain if it's feasible on a single preemptive processor.

## Deadline Monotonic and Rate Monotonic Analysis

The RM priority ordering is not very good when we have tasks with deadline smaller than period, ( $D < T$ ): an infrequent but periodic task would still be given a low priority (because  $T$  is large) and hence probably miss its deadline. For a set of periodic independent tasks, with deadlines within the period, the optimum priority assignment is the *Deadline Monotonic* (DM) priority ordering, which has the following characteristics:

- Priorities are assigned according to task deadlines.
- A task with shorter deadline is assigned a higher priority.
- Tasks with higher priority can preempt lower priority tasks.

The RM priority ordering can be seen as a special case of DM priority ordering. For tasks scheduled by the DM priority ordering there does not exist any simple test for schedulability, like Equation 2 on page 3, but instead we have to calculate the worst case response time for each task separately and compare it to its deadline.

The worst possible response time for a task with any fixed priority assignment policy, (having  $D_i \leq T_i$ ), e.g. DM or RM priority ordering, can be calculated using a response time formula:

$$R_i = C_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j \quad (3)$$

where  $\lceil x \rceil$  is the ceiling operator and calculates the smallest integer  $\geq x$ . Note that the response time variable  $R_i$  is present on both sides of the equation.

The response time analysis builds upon the critical instant theorem and calculates the time a task will have completed if it is activated at the same time as all other tasks with higher priority. The summation gives us the number of times tasks with higher priority will execute before task  $\tau_i$  has completed. The response time for a task can be solved using iteration:

$$R_i^{n+1} = C_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i^n}{T_j} \right\rceil C_j \quad (4)$$

The iteration starts by assuming that the response time is equal to 0. It continues thereafter until two on each other following values of  $R$  becomes equal or  $R$  becomes larger than its deadline  $D$ .

This kind of fixed point response time analysis is (for historical reasons) called *Rate Monotonic Analysis*.

### Assignment 1 cont.

- 1.7 Insert the task set given in Figure 3 on the preceding page in FpsCalc and calculate the response time of each task. What is the worst-case response time for each task using FpsCalc? Verify that the times correspond to the times you extracted using a critical instant schedule. In the worst case scenario, how many instances of  $\tau_1$  and  $\tau_2$  respectively can appear during one execution of  $\tau_3$ ? How does this value relate to the  $\left\lceil \frac{R_i}{T_j} \right\rceil$  expression? How long time of the worst case response time of  $\tau_3$  is spent waiting for instances of  $\tau_1$  and  $\tau_2$  respectively?

### Assignment 2

Task	$C_i$	$T_i$	$D_i$
$\tau_1$	2 ms	20 ms	6 ms
$\tau_2$	3 ms	7 ms	7 ms
$\tau_3$	5 ms	14 ms	13 ms
$\tau_4$	4 ms	100 ms	60 ms

Figure 4: A number of periodic tasks with  $D_i \leq T_i$

2.1 Given the task set in Figure 4 what is the priority ordering for the tasks using the DM priority ordering? What is the priority ordering using RM priority ordering? Use FpsCalc to calculate the response time for the tasks in both ordering. Will all tasks complete before their deadlines? If you are not convinced by the formulas you can do a critical instant schedule.

2.2 Sometimes it is preferable to not use strict RM or DM priority assignment when giving priorities to tasks. This can for example happen when we want to give a task with low deadline demands a better service rate or when the system is part of a larger distributed system.

Find two different priority assignments of the tasks in Figure 4 which is neither RM or DM and where deadlines are missed and met respectively.

2.3 Assume that we want to implement the tasks given in Figure 4 on a RT-kernel that only supports 3 priority levels and where tasks with the same priority will be handled in FIFO order by the scheduler. Assume that task  $\tau_2$  and  $\tau_3$  are set to have the same priority and that we use a DM priority assignment.

Define how the response time formula in Equation 3 on page 4 will be changed when we allow several tasks to have the same priority. Make sure that it is shown in your formula that a task, due to the FIFO order, might have to wait for one instance, but can't be preempted, of an equal priority task. How will the corresponding FpsCalc formula look like, (observe that `sigma(ep, ...)` includes the current task,  $\tau_i$ )?

What will now the worst case response time for each task be? Will all tasks meet their deadlines? Will the worst case response time for task  $\tau_1$  or  $\tau_4$  be affected? Conclusions?

## Blocking

One of the restrictions with the previous analyses is that no tasks are allowed to block or voluntarily suspend themselves. This means that it is very difficult to share resources, such as data, between tasks. In almost all systems where you have several tasks you also have shared resources, e.g. non-preemptible regions of code, FIFO queues and synchronization primitives. The code where a task is accessing shared resources are called *critical sections*, and must often be protected, e.g. by using semaphores.

The problem with using protection mechanisms, such as semaphores, is that high priority tasks might get blocked by lower priority tasks while waiting for a semaphore to be released.

### Assignment 3

3.1 Assume that we have the tasks shown in Figure 6 on the next page. Give priorities to the tasks according to the DM priority assignment. Will all tasks always meet their deadlines?

3.2 Assume that task  $\tau_2$  and  $\tau_4$  in Figure 6 are sharing a semaphore  $S_1$  and that  $\tau_2$  and  $\tau_4$  executes for at most 1 ms and 2 ms respectively in the critical section. Show, by doing a critical instant scheme that the deadline for task  $\tau_2$  can be missed if semaphores and no mechanism for limiting priority inversion (see below) is used. Tip: You might have to model that  $\tau_4$  has taken the semaphore, just before the moment where you let all the other tasks start executing at the same time.

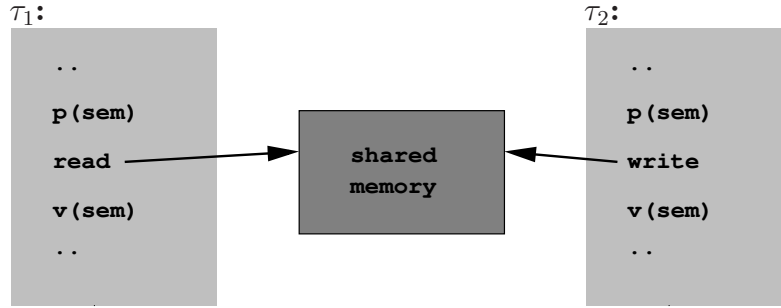


Figure 5: Reading and writing of shared memory using semaphores

Task	$C_i$	$T_i$	$D_i$
$\tau_1$	2 ms	10 ms	5 ms
$\tau_2$	3 ms	20 ms	12 ms
$\tau_3$	10 ms	40 ms	40 ms
$\tau_4$	4 ms	100 ms	50 ms

Figure 6: Four example tasks

The phenomena that a higher priority task not only can be blocked by tasks accessing the same resource but also be delayed by tasks with priorities between the blocking task and the higher priority task is called *priority inversion*.

One protocol that handles the priority inversion is *the priority inheritance protocol*. It works by letting a low priority task that has a semaphore inherit the priority of the highest priority task blocked waiting for the semaphore. When the lower priority task releases the semaphore it will go back to the priority it had before. Notice that the priority of a task now is dynamic so the name 'Fixed Priority' is not really true anymore.

### Assignment 3 cont.

Semaphore	Accessed by	Time
$S_1$	$\tau_2$	1
$S_1$	$\tau_4$	2
$S_2$	$\tau_2$	1
$S_2$	$\tau_3$	5

Figure 7: Tasks semaphore usage

- 3.3 Assume that task  $\tau_2$  not only is sharing a semaphore  $S_1$  with  $\tau_4$  but also is sharing a semaphore  $S_2$  with task  $\tau_3$  as given in Figure 7. Also assume that the times the task  $\tau_2$  is accessing the semaphores  $S_1$  and  $S_2$  do not overlap. Show, by doing a critical instant scheme that deadline for task  $\tau_2$  can be missed even though the priority inheritance protocol is used.

If we want to add the cost for blocking,  $B_i$ , in our response time formula we can do it like:

$$R_i = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil C_j$$

For a task  $\tau_i$ , let  $m$  be the sum of the number of semaphores accessed by  $\tau_i$  and the number of semaphores shared by a lower priority task (than  $\tau_i$ ) and a higher priority task. Using the priority

inheritance protocol, it can be proved that  $\tau_i$  can be blocked at most  $m$  times. It can also be proved that if we are using the priority inheritance protocol  $\tau_i$  can be blocked at most once per lower priority task, and only during one critical section per semaphore. Secondly, for a lower priority task to block  $\tau_i$  it must access a semaphore which a task with higher or equal priority to  $\tau_i$  also is accessing. It gets quite complicated to calculate an exact blocking value for the priority inheritance protocol, mainly due to that all combinations of lower priority tasks and semaphores must be tried, (we get a problem with exponential complexity).

Instead, a simpler, (but somewhat pessimistic), algorithm for calculating the blocking  $B_i$  for  $\tau_i$  can be used: 1. for each semaphore accessed both by a task with priority lower than  $\tau_i$  and by a task with higher or equal priority than  $\tau_i$  we extract the maximum time any lower priority task is accessing the semaphore. 2. set  $B_i$  to be the sum of the  $m$  largest timing values extracted in step 1. The overestimation we get is that the same lower priority task might cause two or more max terms in the summation.

### Assignment 3 cont.

- 3.4 What are the blocking times for the tasks in Figure 6 using the priority inheritance protocol? Make sure that the amount of blocking corresponds to your drawn schedule. Tip: several tasks will experience blocking. What are the response times for the tasks when using the priority inheritance protocol?

Schemes that try to minimize the blocking time of high priority tasks is the *priority ceiling protocol* and the *immediate inheritance protocol*. Since they have the same worst case blocking time we will only investigate the immediate inheritance protocol.

The immediate inheritance protocol works by, in advance for each semaphore, calculate the highest priority of all tasks that will lock the semaphore. This is called the *priority ceiling* for the semaphore. When a task is locking the semaphore, its priority will *immediately* be raised to the maximum of its current priority and the priority ceiling of the semaphore. When the task leaves the semaphore its priority will go back to what it was before.

It can be proven, when using the priority ceiling protocol or the immediate inheritance protocol, that a higher priority task  $\tau_i$  can be blocked at *most once* by a lower priority task. The blocking time is at most *one* critical section of any lower priority task locking a semaphore with ceiling greater than or equal to the priority of task  $\tau_i$ . The blocking term can be given by:

$$B_i = \max_{\{j,s \mid j \in lp(i) \wedge s \in cs(j) \wedge \text{ceil}(s) \geq P_i\}} CS_{j,s}$$

Short explanation: 1. Look at all tasks with priority lower than  $\tau_i$ . 2. Look at all semaphores that these tasks can lock and single out the semaphores where the ceiling of the semaphore is higher than or equal to the priority of  $\tau_i$ . 3. For each semaphore extracted in step 2. extract the largest computation time that the semaphore is locked by one task with priority lower than  $\tau_i$ , ( $CS_{j,s}$  is the computational time for task  $\tau_j$  to execute in the  $s$ :th critical section). 4. Set the largest computation time extracted in step 3. to the blocking time  $B_i$  of task  $\tau_i$ .

Observe that you can use FpsCalc to double-check your results.

### Assignment 3 cont.

- 3.5 What are the priority ceilings for the semaphores  $S_1$  and  $S_2$ ? What are the blocking times for the tasks in Figure 6 using the immediate inheritance protocol? Tip: several tasks will experience blocking. Explain why task  $\tau_3$  can experience blocking even though it does not share any semaphore with  $\tau_4$ . Will all tasks complete before their deadlines?

## Jitter

In the simple model, all processes are assumed to be periodic and to be released with perfect periodicity. This is not, however, always a realistic assumption. *Jitter* is the difference between the earliest and latest start of the task execution relative to the periodic arrival of the task. I.e. it is not certain that a task can start executing with a precise periodicity even if there is no higher priority tasks executing at the same time. The jitter  $J_i^k$  for a task instance  $\tau_i^k$  of a task type  $\tau_i$  is the difference between the arrival time of  $\tau_i^k$  and the time when it gets released. Let  $J_i^{max}$  denote the maximal jitter for the task type  $\tau_i$  (over all possible instances in a given system),  $J_i^{min} = \max\{J_i^k | \forall k\}$ . Let  $J_i^{min}$  denote the minimal jitter for the task type  $\tau_i$  (over all possible instances in a given system),  $J_i^{min} = \min\{J_i^k | \forall k\}$ . The jitter  $J_i$  for a task  $\tau_i$  is the difference between the maximal jitter and the minimal jitter,  $J_i = J_i^{max} - J_i^{min}$ . Figure 8 illustrates how the jitter of a task instance relates to some other already known concepts.

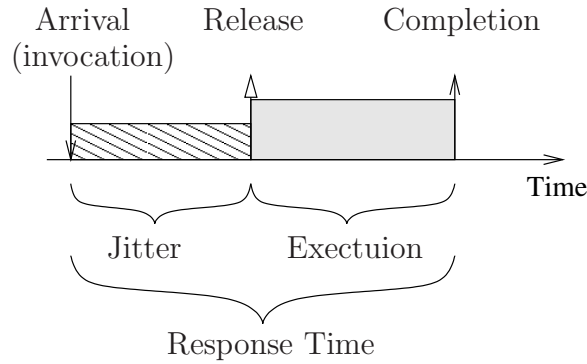


Figure 8: Arrival, release, start of the execution, and jitter of one concrete task instance. The release time and the start of the execution coincide in this example.

One source of jitter is a tick-driven scheduler, i.e. we can only do context switches at regular intervals invoked by the timer hardware. For example see Figure 9 where a tick-driven scheduler is invoked, by a timer-interrupt, every 10th ms, to see if any new task instance has arrived. In case Figure 9(a) we are lucky, since  $\tau_i$ 's instance arrives just before the scheduler got invoked and  $\tau_i$  will start to execute immediately after. In case Figure 9(b) we are more unlucky since  $\tau_i$ 's instance arrives just after the last instance of the scheduler. We now have to wait 10 more ms before the scheduler gets invoked and performs a context switch to task  $\tau_i$ . The jitter of task  $\tau_i$  will therefore be 10 ms, ( $|0 - 10| = 10$ ), due to the tick-driven scheduler.

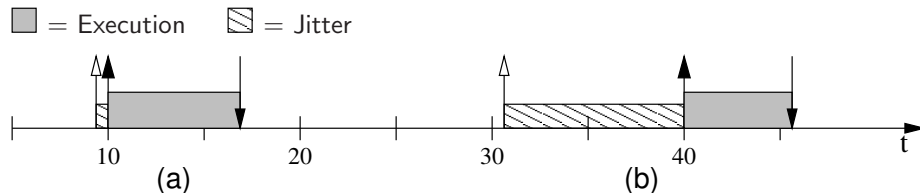


Figure 9: Release jitter due to tick scheduling every 10ms

## Assignment 4

- 4.1 Another source of jitter is varying execution and response times for tasks (or messages) that start other tasks. To illustrate this assume a system with two tasks  $\tau_1$  and  $\tau_2$ . Let task  $\tau_1$  have a period  $T_1 = 10$  and a fixed (non-varying) execution time  $C_1 = 3$ . Let task  $\tau_2$  arrives always 2 time units



after  $\tau_1$  arrives, but it waits for  $\tau_1$  finishing before it gets released (we can see it as that  $\tau_2$  waits for the result of  $\tau_1$  before it can execute). Assume task  $\tau_1$  always ends its execution by releasing task  $\tau_2$  ( $\tau_2$  can now be scheduled for execution). Let  $\tau_2$  have a worst case execution time of  $C_2 = 2$ .

Draw a schedule that shows two instances of  $\tau_1$  and  $\tau_2$  respectively. What is the period  $T_2$  of task  $\tau_2$ ?

- 4.2 To illustrate that varying execution time of  $\tau_1$  might cause jitter of  $\tau_2$  assume that  $\tau_1$ 's execution time no longer is fixed but varies between  $C_1^{min} = 3$  and  $C_1^{max} = 5$ . Task  $\tau_1$  still starts task  $\tau_2$  at the end of its execution.

Draw a schedule with two instances of  $\tau_1$  that shows that the varying execution time of  $\tau_1$  might give raise to jitter of  $\tau_2$ . What is the jitter that task  $\tau_2$  can experience?

- 4.3 To illustrate that interference of high priority tasks might give raise to further jitter of low priority tasks we add a task  $\tau_0$  to the system. Let task  $\tau_0$  have higher priority than both  $\tau_1$  and  $\tau_2$ , a period  $T_0 = 20$  and a worst case execution time  $C_0 = 2$ .

Draw a schedule that shows that varying response time of  $\tau_1$  due to interference of  $\tau_0$  will give raise to further jitter of  $\tau_2$ . Task  $\tau_1$ 's execution time still varies between  $C_1^{min}$  and  $C_1^{max}$ . What is the jitter that  $\tau_2$  can experience due to varying response and execution time of  $\tau_1$ ?

After the above example of how jitter can be derived we will now investigate how jitter of different tasks might interact to cause larger response times of lower priority tasks.

#### Assignment 4 cont.

Task	$C_i$	$T_i$	$D_i$	$J_i$
$\tau_A$	5 ms	20 ms	10 ms	5 ms
$\tau_B$	30 ms	50 ms	50 ms	10 ms

Figure 10: Two example tasks with jitter

- 4.4 For the given tasks  $\tau_A$  and  $\tau_B$  in Figure 10 with DM priority ordering, what is the worst case response time for respective task assuming that we have no jitter. Will both tasks be able to complete before their deadlines?

Assuming that the tasks are experiencing jitter as given in Figure 10 we can use the formula given in Equation 5 to calculate how this will affect the worst case response time of the tasks.

$$\begin{aligned}
 w_i &= C_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i + J_j}{T_j} \right\rceil C_j \\
 &= C_i + \sum_{j \in hp(i)} \left( 1 + \left\lceil \frac{w_i - (T_j - J_j)}{T_j} \right\rceil \right) C_j \\
 R_i &= w_i + J_i
 \end{aligned} \tag{5}$$

#### Assignment 4 cont.

- 4.5 Given the formula in Equation 5 what is the worst case response time for respective tasks assuming that they can experience jitter? Will both tasks be able to complete before their deadlines?

We can construct a task schedule giving the same worst case response time for tasks as the formula in Equation 5.

To construct it we can observe that the critical instant is when both tasks *start executing* at the same time (i.e. not when they arrive at the same time) *and* both tasks experience maximum jitter for their first, (but not the following), instances. As illustrated in Figure 8 the response time of a

task will be the time from which it wants to start executing (arrival time) and the time when it completes its execution. When a task experiences jitter another task can of course execute (if it wants to).

Looking in more detail at the formula in Equation 5 you can see that it is divided into two parts. The  $R_i = \dots$  part express that task  $\tau_i$  can not get preempted when it experiences jitter. The jitter  $J_i$  must still be added to the response time for task  $\tau_i$  (as illustrated in Figure 8).

The  $w_i = \dots$  part represent the time during which  $\tau_i$  can get preempted by higher priority tasks. The  $(1 + \lceil \frac{w_i - (T_j - J_j)}{T_j} \rceil)$  expression captures the release scheme of a higher priority task  $\tau_j$  in which it will be able to preempt  $\tau_i$  as much as possible. This is, in the worst case the first instance of the  $\tau_j$  will have maximum jitter, while the rest of the instances will have minimum jitter.

#### **Assignment 4 cont.**

- 4.6 *Draw a task schedule for the tasks  $\tau_A$  and  $\tau_B$  in Figure 10 which gives the same worst case response times as the formula in Equation 5. Indicate arrival, jitter, beginning of execution, execution, preemption, and completion for each task in your schedule.*

## Report

A proper electronic report is to be handed in for the assignment. It should include at least the following:

- Answers to the assignments. When asked for, illustrate your answers by drawing timing schedules.
- Informative listings of test runs in FpsCalc.

The assignment must be handed in no later than Sunday , 12th Dec., 23:59am.

Some general guidelines of how to make a report:

- Use a word processor, text editor or typewriter to type your solutions. Pictures and diagrams may be drawn by hand and inserted to your editor as picture.
- Provide answers to the questions. Answers should be extracted from FpsCalc files and results. (I.e. do not give answers like “see foo.fps for equation”.)